

Program verification under weak memory consistency

Viktor Vafeiadis

MPI-SWS

Marktoberdorf, August 2019

Axiomatic memory models


Executions

Events

- Reads, Writes, Updates, Fences

Relations

- Program order, po (also called “sequenced-before”, sb)
- Reads-from, rf



Definition (Label)

A *label* has one of the following forms:

$$R \ x \ v_r \quad W \ x \ v_w \quad U(x \ v_r \ v_w) \quad F$$

where $x \in \text{Loc}$ and $v_r, v_w \in \text{Val}$.

Definition (Event)

An *event* is a triple $\langle id, i, l \rangle$ where

- ▶ $id \in \mathbb{N}$ is an event identifier,
- ▶ $i \in \text{Tid} \cup \{0\}$ is a thread identifier, and
- ▶ l is a label.

Definition (Execution graph)

An *execution graph* is a tuple $\langle E, po, rf \rangle$ where:

- ▶ E is a finite set of events
- ▶ po (“*program order*”) is a partial order on E
- ▶ rf (“*reads-from*”) is a binary relation on E such that:
 - ▶ For every $\langle w, r \rangle \in rf$
 - ▶ $typ(w) \in \{W, U\}$
 - ▶ $typ(r) \in \{R, U\}$
 - ▶ $loc(w) = loc(r)$
 - ▶ $val_w(w) = val_r(r)$
 - ▶ rf^{-1} is a function
(that is: if $\langle w_1, r \rangle, \langle w_2, r \rangle \in rf$ then $w_1 = w_2$)




Some notations

Let $G = \langle E, po, rf \rangle$ be an execution graph.

- ▶ $G.E \triangleq E$
- ▶ $G.po \triangleq po$
- ▶ $G.rf \triangleq rf$
- ▶ $G.R \triangleq \{r \in E \mid \text{typ}(r) = R \vee \text{typ}(r) = U\}$
- ▶ $G.W \triangleq \{w \in E \mid \text{typ}(w) = W \vee \text{typ}(w) = U\}$
- ▶ $G.RMW \triangleq \{u \in E \mid \text{typ}(u) = U\}$
- ▶ $G.F \triangleq \{f \in E \mid \text{typ}(f) = F\}$
- ▶ $G.R_x \triangleq G.R \cap \{r \in E \mid \text{loc}(r) = x\}$
- ▶ ...

Mapping programs to executions: Example

Store buffering (SB)

$$\begin{array}{l} x = y = 0 \\ x := 1 \parallel y := 1 \\ a := y \parallel b := x \end{array}$$


Consistency predicate

Let X be some consistency predicate (on execution graphs)

Definition (Allowed outcome under a declarative model)

An outcome O is *allowed* for a program P under X if there exists an execution graph G such that:

- ▶ G is an execution graph of P with outcome O .
- ▶ G is X -consistent.

Exception: “catch-fire” semantics

... or if there exists an execution graph G such that:

- ▶ G is an execution graph of P .
- ▶ G is X -consistent.
- ▶ G is “bad”.

Completeness

The most basic consistency condition:

Definition (Completeness)

An execution graph G is called *complete* if

$$\text{codom}(G.\text{rf}) = G.R$$

i.e., *every* read reads from *some* write.

Sequential consistency

the result of any execution is the same as if the operations of all the processors were executed in some sequential order, respecting the order specified by the program

[Lamport, 1979]

Sequential consistency [Lamport]

Definition

Let sc be a total order on $G.E$. G is called *SC-consistent* wrt sc if the following hold:

- ▶ If $\langle a, b \rangle \in G.po$ then $\langle a, b \rangle \in sc$.
- ▶ If $\langle a, b \rangle \in G.rf$ then $\langle a, b \rangle \in sc$ and there does not exist $c \in G.W_{loc(b)}$ such that $\langle a, c \rangle \in sc$ and $\langle c, b \rangle \in sc$.

Definition

An execution graph G is called *SC-consistent* if the following hold:


- ▶ G is complete.
- ▶ G is SC-consistent wrt some total order sc on $G.E$.

SB example


Store buffering (SB)

$$\begin{array}{l} x = y = 0 \\ x := 1 \parallel y := 1 \\ a := y \parallel b := x \end{array}$$

Allowed



Forbidden



Sequential consistency (Alternative)

Definition (Modification order (aka coherence order))

mo is called a *modification order* for an execution graph G if $mo = \bigcup_{x \in Loc} mo_x$ where each mo_x is a total order on $G.W_x$.

Definition (Alternative SC definition)

An execution graph G is called *SC-consistent* if the following hold:


- ▶ G is complete
- ▶ There exists a modification order mo for G such that $G.po \cup G.rf \cup mo \cup rb$ is acyclic where:
 - ▶ $rb \triangleq G.rf^{-1}; mo \setminus id$ (from-reads / reads-before)

SB example


Store buffering (SB)

$$\begin{array}{l} x = y = 0 \\ x := 1 \parallel y := 1 \\ a := y \parallel b := x \end{array}$$

Allowed



Forbidden



Equivalence

Theorem

The two SC definitions are equivalent.

Proof (sketch).

Lamport SC \Rightarrow alternative SC:

- ▶ Take $\text{mo}_x \triangleq [W_x]; \text{sc}; [W_x]$.
- ▶ Then, $G.\text{po} \cup G.\text{rf} \cup \text{mo} \cup \text{rb} \subseteq \text{sc}$.

Alternative SC \Rightarrow Lamport SC:

- ▶ Take sc to be any total order extending $G.\text{po} \cup G.\text{rf} \cup \text{mo} \cup \text{rb}$.



Relaxing sequential consistency

- ▶ SC is very expensive to implement in hardware.
- ▶ It also forbids various optimizations that are sound for sequential code.

What most hardware guarantee and compilers preserve is “SC-per-location” (aka *coherence*).

Definition

An execution graph G is called *coherent* if the following hold:

- ▶ G is complete
- ▶ For every location x , there exists a total order sc_x on all accesses to x such that:
 - ▶ If $\langle a, b \rangle \in [RW_x]; G.po; [RW_x]$ then $\langle a, b \rangle \in sc_x$
 - ▶ If $\langle a, b \rangle \in [W_x]; G.rf; [R_x]$ then $\langle a, b \rangle \in sc_x$ and there does not exist $c \in G.W_x$ such that $\langle a, c \rangle \in sc_x$ and $\langle c, b \rangle \in sc_x$.

Alternative definition of coherence I

SC: $po \cup \text{rf} \cup \text{mo} \cup \text{rb}$ is acyclic

COH: $po|_{\text{loc}} \cup \text{rf} \cup \text{mo} \cup \text{rb}$ is acyclic

Definition


Let mo be a modification order for an execution graph G . G is called *coherent wrt* mo if $G.po|_{\text{loc}} \cup G.\text{rf} \cup \text{mo} \cup \text{rb}$ is acyclic (where $\text{rb} \triangleq G.\text{rf}^{-1}; \text{mo} \setminus \text{id}$).

Theorem

An execution graph G is coherent iff the following hold:

- ▶ *G is complete*
- ▶ *G is coherent wrt some modification order mo for G .*

“Bad patterns” I



$a := x \text{ // } 1$

$x := 1$

no-future-read




$r := \mathbf{CAS}(x, 1, 1) \text{ // } 1$

rmw-1


Recall:

- ▶ W is either a write or an RMW.
- ▶ R is either a read or an RMW.


“Bad patterns” II

$$\begin{array}{l} x := 1 \quad \parallel \quad a := x \text{ // } 2 \\ x := 2 \quad \parallel \quad a := x \text{ // } 1 \end{array}$$



coherence-ww



coherence-rw



coherence-wr




coherence-rr

“Bad patterns” III



rmw-2



atomicity

In coherent executions, an RMW event may only read from its immediate **mo**-predecessor.

Alternative definition of coherence II

Theorem

Let mo be a modification order for an execution graph G .

G is *coherent* wrt mo iff the following hold:

- ▶ $rf; po$ is irreflexive. (no-future-read)
- ▶ $mo; po$ is irreflexive. (coherence-ww)
- ▶ $mo; rf; po$ is irreflexive. (coherence-rw)
- ▶ $rf^{-1}; mo; po$ is irreflexive. (coherence-wr)
- ▶ $rf^{-1}; mo; rf; po$ is irreflexive. (coherence-rr)
- ▶ rf is irreflexive. (rmw-1)
- ▶ $mo; rf$ is irreflexive. (rmw-2)
- ▶ $rf^{-1}; mo; mo$ is irreflexive. (rmw-atomicity)

Examples (aka “litmus tests”)

Coherence test

$$\begin{array}{lcl} & x = 0 & \\ x := 1 & \parallel & x := 2 \\ a := x \text{ // } 2 & \parallel & b := x \text{ // } 1 \end{array}$$

Store buffering

$$\begin{array}{lcl} & x = y = 0 & \\ x := 1 & \parallel & y := 1 \\ a := y \text{ // } 0 & \parallel & b := x \text{ // } 0 \end{array}$$

Parallel increment

$$x = 0$$
$$a := \mathbf{FAA}(x, 1) \parallel b := \mathbf{FAA}(x, 1)$$

Guarantees that $a = 1 \vee b = 1$.

Can we implement locks in this semantics?

Spinlock implementation

lock(l) :

$r := 0;$

while $\neg r$ **do**

$r := \mathbf{CAS}(l, 0, 1)$

unlock(l) :

$l := 0$

Implementing locks?

Under COH, the spinlock implementation does not guarantee mutual exclusion.

Spinlock implementation

lock(*l*) :

r := 0;

while $\neg r$ **do**

r := **CAS**(*l*, 0, 1)

unlock(*l*) :

l := 0

Lock example

lock(*l*)

x := 1

a := *y* //0

unlock(*l*)

lock(*l*)

y := 1

b := *x* //0

unlock(*l*)

Message passing

More generally, COH is often too weak:


```
x = y = 0
x := 42; || a := y;
y := 1   || while ¬a do a := y;
          || b := x //0
```

```
x = y = 0
x := 42; || a := y; //1
y := 1   || b := x //0
```

MP is a common programming idiom.

How can we disallow the weak behavior?

Supporting message passing



Solution:

- ▶ Strengthen the notion of an “observed” write.
- ▶ In other words, make rf -edges “synchronizing.”

Release/acquire (RA) memory model

SC: $po \cup rf \cup mo \cup rb$ is acyclic

COH: $po|_{loc} \cup rf \cup mo \cup rb$ is acyclic

RA: $(po \cup rf)^+|_{loc} \cup mo \cup rb$ is acyclic

Definition

Let mo be a modification order for an execution graph G .

G is called *RA-consistent wrt* mo if $(po \cup rf)^+|_{loc} \cup mo \cup rb$ is acyclic for some modification order mo for G (where $rb \triangleq G.rf^{-1}; mo \setminus id$).

Definition

An execution graph G is *RA-consistent* if the following hold:

- ▶ G is complete
- ▶ G is RA-consistent wrt some modification order mo for G .

Alternative definition of RA consistency

Theorem

Let mo be a modification order for an execution graph G . G is **RA-consistent wrt** mo iff the following hold:

- ▶ $(\text{po} \cup \text{rf})^+$ is irreflexive. (no-future-read)
- ▶ $\text{mo}; (\text{po} \cup \text{rf})^+$ is irreflexive. (coherence-ww)
- ▶ $\text{rf}^{-1}; \text{mo}; (\text{po} \cup \text{rf})^+$ is irreflexive. (coherence-wr)
- ▶ $\text{rf}^{-1}; \text{mo}; \text{mo}$ is irreflexive. (rmw-atomicity)

$$\text{COH} < \text{RA} < \text{SC}$$

- ▶ Revisit the MP idiom:

$x := 42$	\parallel	$a := y$
$y := 1$		while $\neg a$ do $a := y$
		$b := x$ <i>//0</i>

- ▶ We only need the last read of y to synchronize.
- ▶ Idea: introduce *access modes*.

$x :=_{\text{rlx}} 42$	\parallel	$a := y_{\text{rlx}}$
$y :=_{\text{rel}} 1$		while $\neg a$ do $a := y$
		$a := y_{\text{acq}}$
		$b := x_{\text{rlx}}$ <i>//0</i>

Happens-before

Each memory accesses has a *mode*:

- ▶ Reads: **rlx**, **acq**, or **sc**
- ▶ Writes: **rlx**, **rel**, or **sc**
- ▶ RMWs: **rlx**, **acq**, **rel**, **acq-rel**, or **sc**

“Strength” order \sqsubseteq is given by:



Synchronization:

$$G.\text{sw} = [W \sqsupseteq^{\text{rel}}]; G.\text{rf}; [R \sqsupseteq^{\text{acq}}]$$

Happens-before:

$$G.\text{hb} = (G.\text{po} \cup G.\text{sw})^+$$

Towards C/C++11 memory model

SC: $po \cup rf \cup mo \cup rb$ is acyclic

COH: $po|_{loc} \cup rf \cup mo \cup rb$ is acyclic

RA: $(po \cup rf)^+|_{loc} \cup mo \cup rb$ is acyclic

C11: $hb|_{loc} \cup rf \cup mo \cup rb$ is acyclic

Definition

Let mo be a modification order for an execution graph G . G is called *C11-consistent wrt* mo if $hb|_{loc} \cup rf \cup mo \cup rb$ is acyclic (where $rb \triangleq G.rf^{-1}$; $mo \setminus id$).

Definition

An execution graph G is C11-consistent if the following hold:

- ▶ G is complete
- ▶ G is C11-consistent wrt some modification order mo for G .

The C/C++11 memory model

non-atomic \sqsubseteq relaxed \sqsubseteq release/acquire \sqsubseteq sc

The full C/C++11 is more general:

- ▶ Non-atomics for non-racy code (the default!)
- ▶ Four types of fences for fine grained control
- ▶ SC accesses to ensure sequential consistency if needed
- ▶ More elaborate definition of **sw** (“release sequences”)

C11 model through examples

C11 model through examples

1

```
int a = 0;  
int x = 0;  
a = 42; || if(x == 1){  
x = 1;    ||     print(a);  
          || }  
          ||
```

C11 model through examples

1

```
int a = 0;  
int x = 0;  
a = 42; || if(x == 1){  
x = 1;    ← race    print(a);  
          ||      }  
          ||
```

C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    ← race print(a);
          || }
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xrlx == 1){
xrlx = 1; ||      print(a);
          || }
```

C11 model through examples

1

```
int a = 0;  
int x = 0;  
a = 42; || if(x == 1){  
x = 1;    ← race    print(a);  
          || }  
          ||
```

2

```
int a = 0;  
atomic_int x = 0;  
a = 42; || if(x_rlx == 1){  
x_rlx = 1; ← race    print(a);  
          || }  
          ||
```

C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    ← race      print(a);
          || }
          ||
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xrlx == 1){
xrlx = 1; ← race      print(a);
          || }
          ||
```

3

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xacq == 1){
xrel = 1; ||      print(a);
          || }
          ||
```

C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    ← race      print(a);
          || }
          ||
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xrlx == 1){
xrlx = 1; ← race      print(a);
          || }
          ||
```

3

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xacq == 1){
xrel = 1; ← rf      print(a);
          || }
          ||
```

C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    ← race      print(a);
          || }
          ||
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xrlx == 1){
xrlx = 1; ← race      print(a);
          || }
          ||
```

3

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xacq == 1){
xrel = 1; ← rf, sw → print(a);
          || }
          ||
```


C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    race    print(a);
          }
          ||
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(x_rlx == 1){
x_rlx = 1; race    print(a);
          }
          ||
```

3

```
int a = 0;
atomic_int x = 0;
a = 42; || if(x_acq == 1){
x_rel = 1; rf      print(a);
          sw      }
          ||
```

4

```
int a = 0;
atomic_int x = 0;
a = 42; || if(x_rlx == 1){
fence_rel; fence_acq;
x_rlx = 1;    print(a);
          }
          ||
```

C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    race    print(a);
          }
          ||
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(x_rlx == 1){
x_rlx = 1;    race    print(a);
          }
          ||
```

3

```
int a = 0;
atomic_int x = 0;
a = 42; || if(x_acq == 1){
x_rel = 1;    rf    print(a);
          }
          ||
          sw
```

4

```
int a = 0;
atomic_int x = 0;
a = 42; || if(x_rlx == 1){
fence_rel;    rf    fence_acq;
x_rlx = 1;    print(a);
          }
          ||
```

C11 model through examples

1

```
int a = 0;
int x = 0;
a = 42; || if(x == 1){
x = 1;    ← race    print(a);
           ||      }
```

2

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xrlx == 1){
xrlx = 1; ← race    print(a);
           ||      }
```


3

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xacq == 1){
xrel = 1; ← rf      print(a);
           ||      }
           ||      sw
```

4

```
int a = 0;
atomic_int x = 0;
a = 42; || if(xrlx == 1){
fencerel; ← rf      fenceacq;
xrlx = 1; ← sw      print(a);
           ||      }
```

The “synchronizes-with” relation




$$sw \triangleq ([W^{\sqsupset rel}] \cup [F^{\sqsupset rel}]; po); rf; ([R^{\sqsupset acq}] \cup po; [F^{\sqsupset acq}])$$

Fence modes



Release sequences (RMW's)

$x_{\text{rlx}} := 42;$ \parallel $a := \mathbf{FAI}_{\text{rlx}}(y); \text{ // } 1$ \parallel $b := y_{\text{acq}}; \text{ // } 2$
 $y_{\text{rel}} := 1$ \parallel $c := x_{\text{rlx}}; \text{ // } 0$



$$\text{sw} \triangleq ([W^{\text{rel}}] \cup [F^{\text{rel}}]; \text{po}); \text{rf}^+; ([R^{\text{acq}}] \cup \text{po}; [F^{\text{acq}}])$$

“Catch-fire” semantics

Definition (Race in C11)

Given a C11-execution graph G , we say that two events a, b **C11-race** in G if the following hold:

- ▶ $a \neq b$
- ▶ $\text{loc}(a) = \text{loc}(b)$
- ▶ $\{\text{typ}(a), \text{typ}(b)\} \cap \{W, \text{RMW}\} \neq \emptyset$
- ▶ $\text{na} \in \{\text{mod}(a), \text{mod}(b)\}$
- ▶ $\langle a, b \rangle \notin \text{hb}$ and $\langle b, a \rangle \notin \text{hb}$

G is called **C11-racy** if some a, b C11-race in G .

C11 consistency

Definition

Let mo be a modification order for an execution graph G .
 G is called C11-consistent wrt mo if:

- ▶ $hb|_{loc} \cup rf \cup mo \cup rb$ is acyclic (where $rb \triangleq G.rf^{-1}; mo \setminus id$).
- ▶ ...SC... ?

Definition

An execution graph G is C11-consistent if the following hold:

- ▶ G is complete
- ▶ G is C11-consistent wrt some modification order mo for G .

SC conditions

- ▶ The most involved part of the model, due to the possible mixing of different access modes to the same location.
- ▶ Changed in C++20
- ▶ *If there is no mixing of SC and non-SC accesses*, then additionally require acyclicity of $\text{hb} \cup \text{mo}_{\text{sc}} \cup \text{rb}_{\text{sc}}$.

Further reading:

- ▶ *Overhauling SC atomics in C11 and OpenCL*. Mark Batty, Alastair F. Donaldson, John Wickerson, POPL 2016.
- ▶ *Repairing sequential consistency in C/C++11*. Ori Lahav, Viktor Vafeiadis, Jeehoon Kang, Chung-Kil Hur, Derek Dreyer, PLDI 2017.

Repaired SC condition for fences

$\text{eco} \triangleq (\text{rf} \cup \text{mo} \cup \text{rb})^+ \quad (\text{extended coherence order})$

$\text{psc}_F \triangleq [F^{\text{sc}}]; (\text{hb} \cup \text{hb}; \text{eco}; \text{hb}); [F^{\text{sc}}] \quad (\text{SC fence order})$

Condition on SC fences

psc_F is acyclic

Example: SB with fences

$$\begin{array}{c|c} x = y = 0 & \\ \hline x_{\text{rlx}} := 1; & y_{\text{rlx}} := 1; \\ \text{fence}(\text{sc}); & \text{fence}(\text{sc}); \\ a := y_{\text{rlx}}; \text{ // } 0 & b := x_{\text{rlx}}; \text{ // } 0 \end{array}$$

X behavior disallowed

Exercise: ARC

Add access modes to make the following reference counting implementation correct under C11.

```
new(v){
    a = alloc();
    a.data = v;
    a.count = 1;
    return a;
}

read(a){
    return a.data;
}

clone(a){
    FADD(a.count, +1);
}

drop(a){
    t = FADD(a.count, -1);
    if(t == 1){
        free(a);
    }
}
```

FADD = fetch_and_add